

Palate — an AI marketing studio for a three-venue restaurant brand

The assignment asked me to design two agents on paper. I answered **all four steps and both bonus questions** — and because the best way to judge an agent design is to try it, I also **built the whole thing as a working product**. You can sign in and use it right now.

LIVE DEMO <https://palate-psi.vercel.app>

SIGN IN `owner@marlowandsage.com` · `palate2026`

GITHUB <https://github.com/ArchiusVuong-sudo/palate>

BUILT WITH Next.js 16 · Claude Agent SDK (Opus 4.8) · Gemini 3.1 images · Veo 3.1 video · Supabase

NOTE the hosted demo is great for looking around; for long agent runs and the Chrome setup helper, run it on your own machine (`pnpm install && pnpm dev`) — the hosting plan cuts anything longer than 5 minutes.

Before designing anything — what I would ask the client first


A two-person marketing team doesn't need "an AI strategy". They need hours back, without losing control of how their brand looks and sounds. So before proposing any agents, I would spend an hour with them and ask:

- **Where does the daily hour actually go?** Reading the reviews, or deciding what to do about them? The fix is different for each.
- **What does "off-brand" mean to you?** Show me three posts you rejected and why. Those reasons become the checklist the review agent uses.
- **What should the AI never do on its own?** My starting position: nothing a customer can see goes out without a human clicking "approve".
- **Who owns the accounts?** Do we have admin access to Google Business Profile and Meta, or do we need to sort that out first?
- **What would "this worked" look like in 30 days?** My target: the morning check drops from 60 minutes to a 2-minute read, and the back-and-forth with the designer disappears for everyday posts.
- **What is sacred?** For restaurants: food-safety complaints are never used for marketing; staff are only named with their consent; dietary claims ("gluten free") must be literally true.

The product in this report bakes those answers in as defaults: a human approval step on everything customers can see, brand rules kept as plain-English notes the team can read and edit themselves, and hard limits around safety, consent and claims.

How it's built: one agent, four jobs, many small tools

Instead of four separate black boxes, Palate is **one AI agent** (Claude Opus 4.8) that can be handed four different **jobs** — listening, briefing, creating, reviewing. It works with **23 small tools**, each doing one simple thing: look something up in the database, save an insight, generate an image, ask the team for a decision. All four jobs read the same **brand-knowledge folder** — short plain-English notes about voice, style and rules — and write lessons back into it. Why this matters for the client: when their needs change, we change a few sentences of instructions. Nothing gets rebuilt.



The home screen: today's numbers, the latest creative, the sentiment chart, the agent's suggestions — and a live feed of what the agent is doing right now.

Step 1 — Social Listening & Data Collection

Built as the “Listening” page plus a daily “morning sweep” job. Screenshot below.

1.1 What data sources would this agent monitor?

- **Google Reviews** for each of the three locations — the highest-stakes source: public star ratings, and where food-safety complaints show up.
- **Facebook**: comments on the page’s posts, posts by visitors, and mentions of the page.
- **Instagram**: comments, @-mentions, photo tags and story mentions of the brand account.
- **Worth adding later** (not in the demo): TripAdvisor and OpenTable reviews, plus a simple keyword watch (brand name + dish names) on TikTok — that’s where food trends in Australia usually start.

1.2 What tools or APIs collect the data from each source?

SOURCE	TOOL / API	IN PLAIN TERMS
Google Reviews	Google Business Profile API	Google’s official, free connection. The team grants access once with their Google account; after that, new reviews flow in automatically.
Facebook	Meta Graph API	One Meta developer app covers both Facebook and Instagram. New comments can be pushed to us the moment they appear.
Instagram	Instagram Graph API (same Meta app)	Needs a business Instagram account linked to the Facebook page — most restaurants already have this.
Understanding the text	Claude (the agent itself)	The agent reads each item and tags the mood, the topics and the dishes mentioned. No extra “sentiment software” to buy.

In the demo, the three sources run on realistic made-up data. Palate also has a **setup helper**: the agent opens a real Chrome window, walks through the Google/Meta settings pages with you, you do the sign-ins, and it finds and saves the access keys. For a two-person team, *getting connected* is the real hurdle — not choosing the API.

1.3 What does the agent’s output look like?

One morning run (about 4 minutes, fully automatic) produces three things:

1. **A 30-second digest**, always in the same order: anything urgent first, then trends with numbers, then opportunities, then good news, then one line of stats (mentions in 72h, average mood, % negative, average rating).
2. **Saved insights** — each with a title, an urgency level, the numbers behind it, and *the customers’ own words as quotes*, so the team never has to take a vague claim on trust.
3. **Live charts** the agent publishes itself: mood by day, top topics, which dishes are trending, a 12-week activity map, and a venue-by-venue comparison.

A real example from a run on the demo data: *“CRITICAL (private, never post): Fitzroy served undercooked roast chicken again (9 Jun, 1 star — ‘pink and cold in the middle’). Same problem the*

kitchen was retrained on — it's happening again. Flagged for a same-day temperature-check; keep this off social media entirely."


1.4 How does it tell "important" from noise?

A few simple rules, written into its instructions in English:

- **Safety first, always** — anything about food safety, allergies or staff conduct is urgent and private. No exceptions.
- **Look for change, not volume** — it compares this week against the last one. A dish going from 0 to 3 mentions matters; the 40th compliment for a dish we already know people love doesn't.
- **Every insight must point to an action** — "sentiment is mixed" is not allowed as a finding.
- **Quotes required** — an insight can't be saved without the customer quotes that back it up, so the agent can't make things up.
- **A daily cap** — at most 3–5 insights per day, which forces it to pick what matters most.

1.5 How I'd explain its value in one sentence

"Before you finish your first coffee, it has read every review and comment across all three venues and hands you a 30-second summary — leading with anything that could hurt you, backed by your customers' own words."



Listening: the 14-day mood chart, a 12-week activity map, venue comparison, the analysed feed, and insights with customer quotes.

Step 2 — Daily Content Briefing

2.1 What inputs does this agent take?

- The last 3 days of **insights** from Step 1, with their quotes and numbers, plus any suggestions still waiting.
- The **brand-knowledge notes**: voice, visual style, hard rules, customer types, and facts about each venue (what's on the menu where, which posting times work).
- The **last 5 briefs** — so it doesn't repeat yesterday's idea or feature the same venue twice in a row.
- **Lessons learned** — a notes file the agent adds to over time (for example: "short videos reach about 3× more people than still images").
- An optional **steer from the team** — type "focus on the new winter menu" and the brief follows it.

2.2 What does a good brief look like for this team?

Good enough that the designer doesn't need to ask a single follow-up question:

FIELD	WHY IT STOPS THE BACK-AND-FORTH
Title + date + venue	One brief per day, one clear subject.
Goal	A business outcome ("turn the truffle-gnocchi buzz into bookings"), not just "make a post".
Who it's for	A named customer type ("Date-night Dana") and what she responds to.
Key message	Written the way a customer would say it — usable in the caption as-is.
Tone	Three words from the brand voice notes ("warm, confident, quietly witty").
Channels & formats	Instagram story 9:16 / feed 4:5 / Facebook — decided, not left open.
Visual direction	Points at the style notes: the light, the angle, the colours, what's in the shot.
Wording rules for this post	Must-say and must-not-say ("gnocchi contains gluten — never suggest otherwise").
Call to action + best time to post	What the reader should do, and when the audience is actually looking.
Which insights it came from	So anyone can see <i>why</i> this brief exists.

2.3 An example of the instructions that produce the brief (shortened from the real product)

You are **Palate**, the marketing agent for {brand}. Your job today: write the content brief.

1. Read the latest insights (last 3 days, with quotes) and the open suggestions.
Read the lessons-learned notes and the last 5 briefs, so you never repeat an idea.
2. Pick today's single best idea. Note in one line why it beats the second-best idea.
3. Save ONE brief — so complete the designer asks zero questions: a clear goal, the customer type, the key message in customer language, tone words from the voice notes, channels, call to action, a visual direction that follows our style notes, any must-say / must-not-say wording, and the best time to post.
4. Rules of thumb: back what's rising, not what's already full; never build posts

on safety complaints; claims like "sells out nightly" need a manager to confirm first; vary the venue and the angle across the week.

5. Finish by naming the second-best idea and when it would be the right choice.

A few deliberate choices: *one* brief only (forces a decision), "zero questions" as the quality bar, and the "second-best idea" rule — it makes the agent show its thinking. All the rules are plain sentences the client can change.

2.4 What business logic does it apply?

- **Back what's rising** — a dish going from 0 to 3 mentions with an influencer post beats the dessert that already sells out every night. Posting about the sold-out dish mostly promotes something people can't get.
- **Respond or promote?** — a trending *complaint* (slow Friday service) gets review replies and a note to the kitchen, not a post. You can't post your way out of an operations problem.
- **Avoid friction** — if an idea needs a manager's sign-off (a "selling out fast" claim) and another idea is just as good without it, take the easier one and save the other for tomorrow.
- **The safety wall** — urgent safety issues can never become content. The briefing job simply doesn't see them.
- **Variety** — don't feature the same venue or the same type of idea two days running.

In a real run, the agent picked the truffle-gnocchi launch over the pavlova for exactly these reasons — and wrote the pavlova up as the second-best idea, ready for the day the manager confirms the claim.

2.5 "The brief doesn't feel right for our brand" — what do I adjust?

I treat that sentence as **a lesson to capture, not a complaint to argue with**:

1. **Find where it missed** — is it the tone? the idea? the venue? the look? Each one maps to a different brand note.
2. **Capture the rule in their words** — the client just tells the chat: "never use the word 'pillowy', and keep feed captions under 60 words." The agent writes both into the voice notes as a new version — you can see exactly what changed, and undo it.
3. **Prove it stuck** — this actually happened during the build: the day after learning those two rules, the agent *caught an already-approved caption breaking them* and flagged it for a rewrite, without being asked.
4. **If it keeps missing** — the team writes feedback on the brief and it's redone against that feedback. If it misses repeatedly, the voice notes are too vague — so we sit down for half an hour and rewrite them together.

Step 3 — Visual Content Generation

3.1 What tools, and why these for a restaurant brand?

- **Gemini image generation** (the newest image model, with a stronger one available when needed) — for three reasons that matter for food: (a) it renders food that looks real, avoiding the fake plasticky look AI food often has; (b) it accepts **real photos as references** — we keep a small library of photos of the actual dishes, so the generated picture matches what a guest is really served. For a restaurant that's about honesty, not just looks; (c) it reliably produces the story (9:16) and feed (4:5, 1:1) shapes we need, at about **13 cents per image**.
- **Vevo for short video** — 6-second vertical clips with sound (steam rising, truffle being shaved), built from the chosen photo so the story and the feed post look like a matching set. Short video is where the free reach is.
- **Captions are written first**, before any images — partly because text printed on images comes out better when the wording is final, and partly because the caption *is* the brand's voice. The picture supports it.

3.2 From brief to image prompt — a real example from the build

The brief said (shortened): *Winter truffle gnocchi — “the Burleigh dish nobody knows about yet”. Instagram story + feed. Make the dish the hero: 45° angle, steam, glossy sauce, timber terrace in the background, golden-hour warmth, terracotta and cream tones.*

The image prompt the agent wrote (real):

45° hero shot of hand-rolled potato gnocchi in a cream ceramic bowl, glossy brown-butter sauce, fresh black winter truffle being shaved over the top mid-shave (hand and shaver only, no posed face), visible steam, shallow depth of field with the bowl tack-sharp; weathered timber terrace table softly blurred behind; natural golden-hour warmth; terracotta and cream accents, eucalyptus-green linen; appetising, editorial food photography, no text.

It also attaches the **real photo** of the dish from the brand library, so the plating and the bowl match reality. For stories, it includes the exact overlay line in quotes (“the one nobody knows about — yet”) plus the brand's wordmark.


3.3 Style rules built into every prompt

- Natural light only, warm golden tone — never studio flash or neon. 45° or top-down angles.
- Brand colours somewhere in the frame (terracotta, cream, eucalyptus green); real textures — timber, linen, ceramic.
- Steam for hot dishes; the background softly blurred; the dish is always the hero of the shot.
- People appear as hands and serving moments only — no posed faces (consent, and the photos don't date).
- Text on images: either “no text”, or the exact wording in quotes. The model never invents its own words.
- Alcohol is only ever shown next to food (a responsible-service rule from the brand's hard rules).

These rules live in one style note the team can open and edit. Change the note, and every future image follows it.

3.4 What does the marketing staff actually receive?

A **ready-to-review package per brief**, appearing live in a gallery as it's made (about 5 minutes, roughly \$1 in total): two creative directions × every needed shape (4–6 images), two caption options per channel with hashtags, and one short video — plus the agent's own recommendation of which caption-and-image pairing to use, and why. The gallery has one-click keep/reject, an **Instagram preview** that shows any candidate inside a phone screen, a story player, and a “refine” box (“make the sauce glossier”) that edits the image in place.



The creative studio: the brand's real dish photos, the gallery of options, captions, and the agent's notes on the right.

3.5 What still needs a human, and why

- **Picking the winner** — taste, and what's actually happening in the venue this week, stay human. The agent recommends; it never decides.
- **The truth check** — does the picture match what guests are served today? Only the staff know.
- **Reading the text on the image** — AI-printed text can come out wrong. Notably, the review agent learned this on its own and started zooming into the image to check the words, instead of trusting the prompt.
- **Real photos still matter** — generated images fill the calendar, but genuine staff and guest photos are what people trust most from a local restaurant. This adds to them; it doesn't replace them.

Step 4 — Brand Consistency Review

4.1 What does a “brand guideline” need to contain — and what do I ask the client to prepare?

Four short documents. We write them together in a workshop — I draft, they correct. (Much easier than asking a busy marketer to write rules from scratch.)

- **Voice notes** — three tone words, spelling and capitalisation rules, emoji policy, a banned-words list (“foodie”, “nom”... later “pillowy”), how many exclamation marks are too many, caption length limits.
- **Style notes** — the visual rules from Step 3, plus how the logo is used.
- **Hard rules** — the never-break ones: staff named only with consent; dietary claims only if literally true; no “selling out!” claims without a manager confirming; alcohol only with food; food-safety topics never on social media.
- **Facts** — customer types and per-venue truths (menus, hours, what’s only available where), so claims can be checked.

4.2 How are the guidelines stored so the agent can use them?

As a handful of short plain-English notes, and the agent reads all of them every time. I deliberately didn’t use a search database for this: the notes are small enough to give the agent in full, so there’s no chance a search step quietly misses the one rule that mattered. The practical wins: *the client can open and edit their own rules*; every change is saved as a new version with a side-by-side diff and an undo; and the agent adds its own learned rules (clearly marked as agent-written) which the team can check. On top of that, the most important rules are **enforced by the software itself**: the publishing tool refuses to run unless a human approval is attached. The instructions ask nicely — the code makes sure.

4.3 What does the agent check? (six things)

1. **Voice & tone** — banned words, emoji and punctuation rules, spelling, caption length.
2. **Look & feel** — light, angle and colours against the style notes — judged by *looking at the image*, not at the prompt that made it.
3. **Hard rules** — consent list, dietary claims, scarcity claims, alcohol-with-food, pricing.
4. **Accuracy** — every factual claim checked against the facts notes (right venue, dish actually on that menu).
5. **Right audience, right format** — the right customer type, and story-style writing for stories, feed-style for the feed.
6. **Text on the image** — the printed words and logo checked visually for spelling and readability.

Each one gets a score out of 100 with a one-line note — so “feels off-brand” becomes a number the team can discuss.


4.4 The three possible outcomes — pass, flag, reject

OUTCOME	WHEN	WHAT HAPPENS
---------	------	--------------

Pass	Everything scores 75+ and no hard rule is broken	The scorecard is saved and the post goes to the human approval step. Passing does not mean published — a person still clicks.
Flag	Anything scores 50–74, or there's a fixable issue	The agent lists exact fixes — and for image problems it <i>draws boxes on the picture</i> showing what's wrong. One retouch is tried, then it's reviewed again.
Reject	A hard rule is broken, or anything scores under 50	A plain-English explanation; the post goes back to “changes requested”; the lesson is added to the notes so the same mistake fades out.


4.5 How do I keep it from being too strict — or too lenient?

- **One simple rule: strict on the hard rules, relaxed on taste.** Only a broken rule can reject a post; matters of taste can only flag it. This one line prevents most wrongful blocking.
- **Numbers, not vibes** — the pass/flag/reject lines are set as numbers, so adjusting strictness means moving a number, not rewriting the agent.
- **The human step is the safety valve** — a too-strict reviewer can’t actually block anything, because the approval card shows the scorecard and the human gets the final say.
- **Disagreements become lessons** — every time the team overrides the agent, that gets written into the notes, so its judgement drifts toward the team’s real taste.
- **A spot check now and then** — re-score a sample of posts the team approved and published. If the agent keeps disagreeing with reality, the notes need a rewrite — not the agent.



The review board: every post moves through the same queue, with scorecards — and a human click as the final step.

Bonus A — How all four agents connect, end to end



What starts each step. Three ways in, same machinery underneath: the 7:00 am *Autopilot* runs the whole cycle daily; any single step can be run from its own page ("Generate today's brief"); and the chat can kick off any step in plain words ("draft a reply to the slow-service review"). Steps hand work to each other through saved, editable pieces — insights → brief → images → post — not through some invisible chain. That means a person can step in anywhere: edit the brief before the images are made, or reject one image without redoing the whole morning.

If a step fails (an outage, a time-out, the spending cap): the run is marked failed and its full step-by-step log is kept. Whatever was already made is saved — if image 4 of 5 failed, you keep the first three. Nothing downstream goes out by itself, because the human step never fired. The team sees the failure in the run log and the daily digest, and a "Continue" button picks the agent up exactly where it stopped, with everything it knew — no starting over. The worst case is always "the team does today the old way", never "something wrong went out".

One real, unplanned example of the whole loop: during the build, the morning Autopilot ran by itself. It read the overnight mentions, noticed that an already-approved caption broke two rules the owner had taught it the day before (the word "pillowy"; a 60-word limit), wrote a corrected caption, and **stopped at the approval step** to ask permission to swap it — quoting the rules by name. Every step, plus the learning loop, working together with nobody pressing a button.

Bonus B — "Can we have one button that does everything automatically, without us checking?"

What I would actually say to the client:

"Almost — and the 'almost' is what protects you. That button exists: Autopilot runs the whole cycle every morning at 7. By the time you open your laptop, the reviews are read, today's brief is written, the visuals are made and checked against your brand. The one thing it won't do is *publish* without you. That last click — about two minutes of your day — stays with you, for three reasons.

"First, the rare cases are the ones that matter. Last week the system caught a review about undercooked chicken. A fully automatic system having a normal day might have posted cheerful food content right on top of that; you'd know instantly not to. You're not checking its spelling — you're adding the judgement

only you have. Second, responsibility: when something goes out under your name, a person chose it. 'The AI posted it' is not an answer you want to give anyone — ever. Third, your two minutes are how it learns: every approve, reject and edit teaches it your taste, so the checking keeps getting faster — it already caught one of its own captions breaking a rule you'd taught it the day before. And it's not just me being careful: every serious AI marketing product on the market keeps a human approval step, for the same reasons. If reviewing ever starts feeling like a chore that isn't shrinking, that's a fault in my system — not a reason to remove the step."


The product in pictures — ten things worth seeing first

Everything below is the real product, captured while it was running. It's in two groups on purpose: first **how the agent works with the team** — because for two overwhelmed people, how it asks, alerts, remembers and helps is worth more than any single output — and then **what it makes**.

Part one — working with the agent

1 · Watch the agent work — and watch it stop to ask you

Every run is a live window you can watch (and replay later): the agent's thinking as it streams, each tool it uses ("Saving review → pass 90/100"), the images and charts it makes — and the moment something customer-facing is at stake, it **stops and asks**. Here the review agent files two scorecards, then says: *"Now I'll request human approval to move the corrected feed post to approved."* Your answer is recorded in the same timeline, the cost is on the receipt (\$1.16), and the "Follow up" box lets you continue the same conversation — even days later, it still remembers everything.



2· Decisions come to you — you never have to go looking

When the agent needs an answer, it makes sure you see it: a glowing “1 decision waiting” button appears in the top bar of every page, the Review menu shows a badge, a small widget sits in the corner, the browser tab title changes, and a notification pops up. The decision card itself carries everything needed to decide on the spot — the image, the caption, the planned time — plus a note box, and whatever you write there becomes a lesson the agent keeps. Built this way for a simple reason: a busy team shouldn’t have to remember to check on their assistant.

The screenshot displays the Palate Marketing Studio interface, which is designed for managing content approvals. The top navigation bar includes a search bar, a "Copilot" button, and a badge indicating "1 decision waiting". The left sidebar lists various studio functions: Overview, Listening, Briefs, Creative, Review (highlighted with a badge), Calendar, Knowledge, Agent runs, Outbox, and Settings.


The main workspace is titled "Brand review" and shows a "PENDING APPROVALS" section with one item. The item is a post scheduled for 6:00 pm today, featuring an image of a dish and a caption about Burleigh's winter menu. The post has passed a brand check with an overall score of 90. Below the image, there is a text box for an optional note and three buttons: "Approve", "Reject", and "Request changes".

Below the pending approvals, the "Publish pipeline" is shown, illustrating the flow from draft to live. The pipeline consists of four stages: IN REVIEW (1 item), CHANGES REQUESTED (0 items), APPROVED (1 item), and SCHEDULED & PUBLISHED (0 items). The "IN REVIEW" stage shows a card for an "IG Feed" post, which has passed. The "APPROVED" stage shows a card for an "IG Story" post, which has also passed.

On the right side, the "Palate copilot" panel provides additional context and instructions. It includes a note about the 60-word caption limit, a reminder to send a marketing digest email, and a list of "Digest covered" items: Listening highlights, Truffle brief, Approved posts, and Critical Fitzroy escalation. It also mentions that the approved IG Feed caption breaks new owner rules and is ~63 words long.

3· It can use a real web browser — and it uses it to help you get set up

The agent doesn't just write and draw — it has **browser abilities**: it can open a real Chrome window on your computer, read what's on a page, click and type, the way a person would. Today that powers the part where non-technical teams usually get stuck: connecting the Google and Meta accounts. Press "Set up with the agent" on any connection card and it drives the settings pages with you watching — and it pauses and hands you the wheel at every sign-in, because it should never see a password. Below: a real run where the person stopped at the sign-in step, and the agent honestly saved the connection as "pending", with a note on exactly where it left off.



4 · A copilot you can simply talk to — and it can look things up

The chat on the right side of every page is the same agent in conversation. Ask “which dish is trending this week?” and it answers with real numbers from your data; tell it “never use the word ‘pillowy’” and it updates the brand notes permanently; ask for a digest email and it drafts one and asks before sending. It remembers earlier conversations, and when a job needs outside information it can also **search the web and read pages** — for example, checking what an influencer actually posted before quoting them in a brief.

Palate

MARKETING STUDIO

Overview

Listening

Briefs

Creative

Review

Calendar

Knowledge

Agent runs

Outbox

Settings

STUDIO / Overview

Search or command

Copilot

Good morning — Friday 12 June

Here's what your customers are saying

MENTIONS · 72H

22 ↗ 83%

vs previous 72 hours

AVG SENTIMENT

0.37

scale -1 to 1 · last 14 days

AVG RATING · 14D

3.80★

out of 5

PENDING APPROVALS

0

waiting on you

Morning sweep

Run full pipeline

Latest creative

Fresh variants from the creative studio

Open studio →

the one nobody knows about — yet

Eye-catching story 9:16

the one nobody knows about — yet

review-markup story 9:16

the one nobody knows about — yet

A truffle shave story 9:16

find it before everyone else

B — table story story 9:16

B — table scene feed 4:5

the one nobody knows about — yet

A — hero story story 9:16

A — feed

Sentiment pulse — 14 days

Daily positive vs negative mentions

12

9

6

Where it's happening

Mentions by source · 14 days

53

MENTIONS

Palate copilot

caption before it goes out. Stories aren't affected by the 60-word cap unless you want me to apply it there too.

Send today's marketing digest email to khanhvuongtuan37@gmail.com — cover the listening highlights, the truffle brief, the approved posts, and the critical Fitzroy escalation. Brand-styled HTML please.

Sent — it's in the Outbox ([e271d9ce](#)) since Gmail SMTP isn't wired up, so it's visible on the Outbox page rather than delivered externally.

Digest covered:

- Listening highlights — pavlova selling out nightly (Surry Hills), truffle buzz (Burleigh), Friday-peak service still slipping, Burleigh parking churn risk
- Truffle brief — "the dish nobody knows about yet" (draft)
- Approved posts — IG Story (ready) + IG Feed (approved, but flagged for edit)
- Critical Fitzroy escalation — framed strictly as internal-only per the food-safety guardrail

One thing that needs your call: the approved IG Feed caption breaks both new owner rules — it contains "pillowy" and is ~63 words. Since it's already in approved status, I haven't touched it. Want me to draft a compliant rewrite (under 60 words, no "pillowy") and put it up for your approval to replace the current one?

Ask, instruct, or correct me...

5 · One click sends anything to the chat — no copy-pasting

Small “Ask Palate” buttons sit on reviews, insights, posts and runs all over the studio. Click one and the chat opens with that item’s full text already typed in — here, a customer review, ready to discuss with one keystroke. You stay in charge (nothing sends until you press enter), but you never have to copy things around or explain the background. The product carries the context for you.

The screenshot displays the Palate Marketing Studio interface, which is divided into several sections:

- Left Sidebar:** Contains navigation links for Overview, Listening (selected), Briefs, Creative, Review, Calendar, Knowledge, Agent runs, Outbox, and Settings.
- Top Bar:** Shows the current view as "STUDIO / Listening" and includes a "1 decision waiting" indicator, a search bar, and a "Copilot" button.
- Main Content Area:**
 - Social listening:** A section titled "Google Reviews · Facebook · Instagram — analysed every morning" with a "Run morning sweep" button.
 - THE PULSE — 14 DAYS:** A line chart showing positive (green) and negative (red) sentiment over time from May 31 to June 12. A peak in positive sentiment is visible around June 10.
 - Summary Cards:** Three cards showing "UNANALYSED 0 in the inbox", "POSITIVE 63% of analysed mentions", and "FLAGGED 1 safety / urgent".
 - Listening rhythm:** A heatmap showing mentions over time, with filters for "54 mentions", "busiest 10 Jun · 14", and "15/84 days active".
- Right Sidebar:**
 - Palate copilot:** A section with a "caption before it goes out" note and a "Send today's marketing digest email" button.
 - Digest covered:** A list of items including "Listening highlights", "Truffle brief", and "Fitzroy escalation".
 - Chat Window:** A pop-up window showing a review from "Paul Stevens" about "Gnocchi" with a "Positive +0.50" score. It includes a "POSTED" date of "11 June, 6:09 pm" and an "ANALYSED" time of "14h ago".

6 · Brand rules with a memory — you can see every change

The brand guidelines are living notes with a history. When the owner typed *“never use the word ‘pillowy’; keep feed captions under 60 words”* into the chat, the agent updated the voice notes itself — the green/red comparison below shows exactly what it added, marked as agent-written, with one click to undo. The next morning it enforced those rules against its own already-approved caption, without being asked. That’s the difference between an instruction and a memory.

The screenshot displays the Palate Marketing Studio interface, specifically the 'Knowledge' section. The left sidebar lists various tools: Overview, Listening, Briefs, Creative, Review, Calendar, Knowledge (selected), Agent runs, Outbox, and Settings. The main content area is titled 'Brand knowledge' and shows a list of files under different categories: BRAND, AUDIENCE, OPERATIONS, and LEARNINGS. The 'voice.md' file is highlighted, showing its history with version 2 marked as 'agent' and version 4 marked as 'agent'. A tip box suggests using the word 'cheap' and saving the lesson. The right panel shows the 'Brand voice' document, which has been updated with two permanent owner directives: 'ban 'pillowy'; feed captions under 60 words.' The document content includes sections for 'Marlow & Sage - Brand Voice', 'Personality', and 'Writing rules'. A comparison view shows the changes made by the agent, with a 'Restore this version' button.

Palate MARKETING STUDIO

STUDIO / Knowledge

Search or command

Copilot

Overview 9

Listening 01

Briefs 02

Creative 03

Review 04

Calendar

Knowledge

Agent runs

Outbox

Settings

Brand knowledge

The agent's memory — it reads this before every job and writes lessons back

+ New file

BRAND 3

guidelines.md v1 human

visual-style.md v1 human

voice.md v2 agent

AUDIENCE 1

personas.md v1 human

OPERATIONS 1

locations.md v1 human

LEARNINGS 2

content-performance... v1 human

feedback-log.md v4 agent

Tip — tell the copilot things like “never use the word cheap” and it will file the lesson here itself.

1 agent run live LISTENING

brand/voice.md · v2

Brand voice

updated 1d ago — “Adding two permanent owner directives to brand voice: ban ‘pillowy’; feed captions under 60 words.”

v2 CURRENT agent “Adding two permanent owner directives to brand voice: ban ‘pillowy’; feed captions under 60 words.” 1d ago

v1 human “Initial brand onboarding” 1d ago

v1 → v2 (current) +3 -0 Restore this version

Marlow & Sage — Brand Voice

One line: Warm, confident, quietly witty. We sound like a great host, not a marketing department.

Personality

- Modern Australian: relaxed, generous, never stuffy or formal.
- Witty but understated — one light touch of humour per post max, never forced slang.
- Proud of the food without bragging; let specifics do the talking (“48-hour lamb shoulder”, not “amazing food”).

Writing rules


- Australian English spelling (flavour, colour, favourite).
- Sentence case for everything, including headlines. No ALL CAPS.
- Maximum one emoji per caption, two on stories. Never use 🍕 or 🍔.
- No exclamation mark pile-ups. One “!” maximum per caption.
- Never use: “foodie”, “nom”, “drool”, “cheat day”, “guilt-free”, “to die for”.
- Hashtags: 3-6, specific over generic (#suiryhillseats over #food). Always include #marlowandsage.
- Refer to staff by first name when celebrating them (with their permission noted).

Preview Close history

Part two — what it makes


8 · The Instagram preview — see the post the way a guest will

Any image-and-caption option can be shown inside a realistic phone screen — story progress bar, account name, “Send message” bar and all. Choosing between option A and option B stops being abstract: the team judges the exact thing their guests will see. There’s also a story player to watch the full sequence, and a before/after slider for retouches.




9 • A content calendar the team and the agent share

Briefs and posts land on a month view with their channel and status visible at a glance. Approved posts can be **dragged to a new day** to reschedule them (saved instantly, 6 pm by default), a small “unscheduled” shelf holds approved work still waiting for a date, and “Ask Palate about gaps” hands the whole month to the agent to plan around.



10 · Cmd-K — the whole studio under one keystroke

One search box finds everything — reviews by customer name, briefs, images with thumbnails, the brand notes — and can also *start agent runs*: type “Run morning sweep” from any page and watch it live in the corner. With keyboard shortcuts, alerts, and an always-visible “decision waiting” button, the studio is quick to drive once you know it.



Beyond the brief — what the working build shows

CAPABILITY	WHY IT MATTERS
An approval step that really blocks	The agent pauses mid-run until a human answers — and publishing or emailing is refused by the software itself without an approval attached. The model can't talk its way past it.
A memory the client owns	Brand notes with versions, comparisons and undo. The agent adds its own lessons, clearly marked, and the team can check them. The "pillowy" rule went from a chat message to an enforced rule overnight.
Conversations that pick up where they left off	The chat remembers earlier sessions, and any finished run can be continued with its full memory — tested across a 19-hour gap.
Agent-assisted setup	The agent walks the team through Google/Meta account connection in a real browser window, with humans doing every sign-in (works when running locally).
A real operations view	Content calendar with drag-to-reschedule, a live feed of agent activity, full run logs with per-run cost, a daily Autopilot switch, search-everything, and an email digest.
Known, measured costs	Measured in the build: morning sweep $\approx$ \$0.76 · brief $\approx$ \$0.40 · creative package $\approx$ \$0.50 + $\sim$ \$0.13 per image + $\sim$ \$0.60 per video · review $\approx$ \$0.98 · chat replies \$0.06–0.29. A full day's cycle is roughly \$2–3 (about A\$100/month) — against 1.5–2 staff hours saved every day.

Being honest about limits — and the first 30 days of real use

- **The demo runs on made-up data.** The real Google/Meta connections plug into the connection cards that already exist (and the setup helper already handles the keys). Weeks 1–2.
- **It doesn't learn from post performance yet** — the next step is feeding actual results (likes, reach, bookings) back in, so briefs and reviews are informed by what worked. The calendar and the scorecards are already built to receive that. Weeks 2–4.
- **Hosted-demo limits** — the hosting plan cuts off long runs at 5 minutes, so the full agent experience is best on your own machine. The production answer is a small separate worker for runs; the code is already organised for that.
- **Spending is capped** — every run has a hard budget limit (\$8) and every tool call and dollar is logged.

5-MINUTE LOOK Sign in at palate-psi.vercel.app (`owner@marlowandsage.com` / `palate2026`).

SUGGESTED

PATH Listening (the activity map, insights with quotes) → Briefs (read “why this angle”) → Creative (gallery, phone preview, the agent’s notes) → Review (the queue + scorecards) → Knowledge (open `brand/voice.md` → History → see the agent’s change) → Agent runs (the logs, costs, and “Continue”).

FULL

EXPERIENCE `git clone https://github.com/ArchiusVuong-sudo/palate && npm i && npm dev` with the provided `.env.local` → press “Run full pipeline” and answer the approval card when the agent asks.

Palate — built end to end (product, agents, deployment, this report) as the AT Solutions take-home by Khanh Vuong, June 2026. All demo data is fictional. Models: Claude Opus 4.8 (the agent), Gemini 3.1 Flash Image, Veo 3.1, via the Claude Agent SDK.